

Long read splice alignment — theory and practice

Kristoffer Sahlin

Department of Mathematics, Science for Life Laboratory,
Stockholm University

Intended Learning Outcomes

Should

- Learn the basics of an aligner (concepts: seeding, chaining, extension)
- Be able to run minimap2 to obtain alignments
- Think critically about reliability of alignments for downstream analysis
- ~~(Learn some methods for basic sanity checking)~~

Workshop overview

Theory

- Long-read alignment (seeding, chaining, extension) à la minimap2
- Long-read splice alignment à la minimap2

Parameters and heuristics

- Seed and window size (k, w) - uniqueness and speed
- Some minimap2 specific parameters
- Thresholds

Exercise

- Mapping transcripts to references with minimap2

More theory

- Aligner variants (uLTRA, deSALT)

Interpreting the output

- SAM format
- MAPQ score and secondary and supplementary alignments
- CIGAR format

~~Troubleshooting~~

- ~~samtools, BLAT, IGV~~

1. Seeding with *k*-mers

- A match between read and reference
- Types of seeds: ***k*-mers** (there are others, e.g., spaced *k*-mers, MEMs etc)
- A *k*-mer is a substring of length *k*

CACGACTCTGGTACCTAGACTCGATCGATCGTACTGT....

CACG

ACGA

CGAC

k-mers with $k=4$

1. Seeding with *k*-mers

- A match between read and reference
- Types of seeds: ***k*-mers** (there are others, e.g., spaced *k*-mers, MEMs etc)
- A *k*-mer is a substring of length *k*

ch1 CACGACTCTGGTACCTAGACTCGATCGATCGTACTGT....
ch2 GCACACTGGTACCTAGACACGATCGCCCGTTGTGTCA....
ch3 GCTGACTGTTACAACCTATACTCGATCGCCCGTTGTCT....
...

Index

k-mer	Positions
CACG	[(ch1,0), (ch17,1202), ...]
AGAC	[(ch1,16), (ch2, 14), ...]
TATA	[(ch2, 25), (ch13, 205), ...]
GACT	...


...

1. Seeding with *k*-mers

- A match between read and reference
- Types of seeds: ***k*-mers** (there are others, e.g., spaced *k*-mers, MEMs etc)
- A *k*-mer is a substring of length *k*

ch1 CACGACTCTGGTACCTAGACTCGATCGTACTGT....
ch2 GCACACTGGTACCTAGACACGATCGCCCGTTGTGTCA....
ch3 GCTGACTGTTACAACCTATACTCGATCGCCCGTTGTCT....
...

?



AGACCCGAT

Read (query)

Index

k-mer	Positions
CACG	[(ch1,0), (ch17,1202), ...]
AGAC	[(ch1,16), (ch2, 14), ...]
TATA	[(ch2, 25), (ch13, 205), ...]
GACT	...
...	

1. Seeding with *k*-mers

- A match between read and reference
- Types of seeds: ***k*-mers** (there are others, e.g., spaced *k*-mers, MEMs etc)
- A *k*-mer is a substring of length *k*

ch1 CACGACTCTGGTACCTAGACTCGATCGATCGTACTGT....
ch2 GCACACTGGTACCTAGACACGATCGCCCGTTGTGTCA....
ch3 GCTGACTGTTACAACCTATACTCGATCGCCCGTTGTCT....
...

?
AGACCCGAT
Read (query)

AGAC


Index

k-mer	Positions
CACG	[(ch1,0), (ch17,1202), ...]
AGAC	[(ch1,16), (ch2, 14), ...]
TATA	[(ch2, 25), (ch13, 205), ...]
GACT	...
...	

1. Seeding with *k*-mers

- A match between read and reference
- Types of seeds: ***k*-mers** (there are others, e.g., spaced *k*-mers, MEMs etc)
- A *k*-mer is a substring of length *k*

ch1 CACGACTCTGGTACCTAGACTCGATCGTACTGT....
ch2 GCACACTGGTACCTAGACACGATCGCCCGTTGTGTCA....
ch3 GCTGACTGTTACAACCTATACTCGATCGCCCGTTGTCT....
...




Index

k-mer	Positions
CACG	[(ch1,0), (ch17,1202), ...]
AGAC	[(ch1,16), (ch2, 14), ...]
TATA	[(ch2, 25), (ch13, 205), ...]
GACT	...
...	

1. Seeding with *k*-mers

- A match between read and reference
- Types of seeds: ***k*-mers** (there are others, e.g., spaced *k*-mers, MEMs etc)
- A *k*-mer is a substring of length *k*

ch1 CACGACTCTGGTACCTAGACTCGATCGTACTGT....
ch2 GCACACTGGTACCTAGACACGATCGCCCGTTGTGTCA....
ch3 GCTGACTGTTACAACCTATACTCGATCGCCCGTTGTCT....
...



Index

k-mer	Positions
CACG	[(ch1,0), (ch17,1202), ...]
AGAC	[(ch1,16), (ch2, 14), ...]
TATA	[(ch2, 25), (ch13, 205), ...]
GACT	...
...	

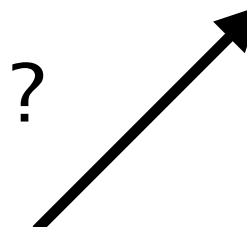
1. Seeding with *k*-mers

- A match between read and reference
- Types of seeds: ***k*-mers** (there are others, e.g., spaced *k*-mers, MEMs etc)
- A *k*-mer is a substring of length *k*

ch1 CACGACTCTGGTACCTAGACTCGATCGTACTGT....
ch2 GCACACTGGTACCTAGACACGATCGCCCCGTTGTGTCA....
ch3 GCTGACTGTTACAACCTATACTCGATCGCCCCGTTGTCT....
...

AGACCCGAT
Read (query)

?



CCCCG




Index

k-mer	Positions
CACG	[(ch1,0), (ch17,1202), ...]
AGAC	[(ch1,16), (ch2, 14), ...]
TATA	[(ch2, 25), (ch13, 205), ...]
GACT	...
...	

1. Seeding with *k*-mers

- A match between read and reference
- Types of seeds: ***k*-mers** (there are others, e.g., spaced *k*-mers, MEMs etc)
- A *k*-mer is a substring of length *k*

ch1 CACGACTCTGGTACCTAGACTCGATCGTACTGT....
ch2 GCACACTGGTACCTAGACACGATCGCCCGTTGTGTCA....
ch3 GCTGACTGTTACAACCTATACTCGATCGCCCGTTGTCT....
...



Index

k-mer	Positions
CACG	[(ch1,0), (ch17,1202), ...]
AGAC	[(ch1,16), (ch2, 14), ...]
TATA	[(ch2, 25), (ch13, 205), ...]
GACT	...
...	

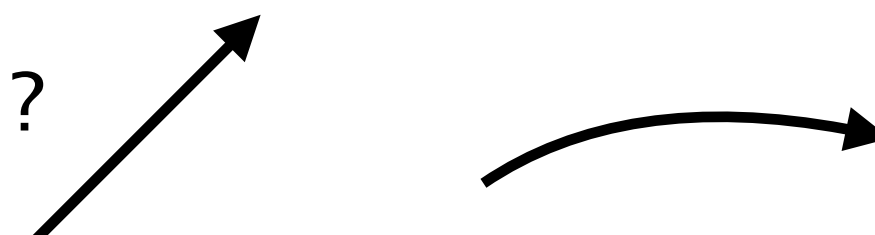
1. Seeding with *k*-mers

- A match between read and reference
- Types of seeds: ***k*-mers** (there are others, e.g., spaced *k*-mers, MEMs etc)
- A *k*-mer is a substring of length *k*

ch1 CACGACTCTGGTACCTAGACTCGATCGATCGTACTGT....
ch2 GCACACTGGTACCTAGACACCGATCGCCCGTTGTGTCA....
ch3 GCTGACTGTTACAACCTATACTCGATCGCCCGTTGTCT....
...

AGACCCGAT
Read (query)

?



CGAT

Index


k-mer	Positions
CACG	[(ch1,0), (ch17,1202), ...]
AGAC	[(ch1,16), (ch2, 14), ...]
TATA	[(ch2, 25), (ch13, 205), ...]
GACT	...
...	

1. Seeding with *k*-mers

- A match between read and reference
- Types of seeds: ***k*-mers** (there are others, e.g., spaced *k*-mers, MEMs etc)
- A *k*-mer is a substring of length *k*

ch1 CACGACTCTGGTACCTAGACTCGATCGATCGTACTGT....
ch2 GCACACTGGTACCTAGACACGATCGCCCGTTGTGTCA....
ch3 GCTGACTGTTACAACCTATACTCGATCGCCCGTTGTCT....
...

?



AGACCCGAT

Read (query)

Index


k-mer	Positions
CACG	[(ch1,0), (ch17,1202), ...]
AGAC	[(ch1,16), (ch2, 14), ...]
TATA	[(ch2, 25), (ch13, 205), ...]
GACT	...
...	

1. Seeding with *k*-mers

- A match between read and reference
- Types of seeds: ***k*-mers** (there are others, e.g., spaced *k*-mers, MEMs etc)
- A *k*-mer is a substring of length *k*

ch1 CACGACTCTGGTACCTAGACTCGATCGATCGTACTGT....
ch2 GCACACTGGTACCTAGACACGATCGCCCGTTGTGTCA....
ch3 GCTGACTGTTACAACCTATACTCGATCGCCCGTTGTCT....
...

?



AGACCCGAT

Read (query)

Index

k-mer	Positions
CACG	[(ch1,0), (ch17,1202), ...]
AGAC	[(ch1,16), (ch2, 14), ...]
TATA	[(ch2, 25), (ch13, 205), ...]
GACT	...
...	

1. Subsampling seeds

- Store only a subset of k-mers (aka *sketching* or *thinning*)
- At least one seed in every “window” - guarantee
- Parameter w (window size) controls density

CACGACTCTGGTACCTAGACTCGATCGATCGTACTGT....

CACG

ACGA

CGAC

GACT

ACTC

“Minimizers”, $k = 4$, $w = 5$

1. Subsampling seeds

- Store only a subset of k-mers (aka *sketching* or *thinning*)
- At least one seed in every “window” - guarantee
- Parameter w (window size) controls density

CACGACTCTGGTACCTAGACTCGATCGATCGTACTGT....

CACG

ACGA

CGAC

GACT

ACTC

CTCT

“Minimizers”, $k = 4$, $w = 5$

1. Subsampling seeds

- Store only a subset of k-mers (aka *sketching* or *thinning*)
- At least one seed in every “window” - guarantee
- Parameter w (window size) controls density

CACGACTCTGGTACCTAGACTCGATCGATCGTACTGT....

CACG

ACGA

CGAC

GACT

ACTC

CTCT

TCTG

“Minimizers”, $k = 4$, $w = 5$

1. Subsampling seeds

- Store only a subset of k-mers (aka *sketching* or *thinning*)
- At least one seed in every “window” - guarantee
- Parameter w (window size) controls density

CACGACTCTGGTACCTAGACTCCATCGATCGTACTGT....


“Minimizers”, $k = 4$, $w = 5$

1. Subsampling seeds

- Store only a subset of k-mers (aka *sketching* or *thinning*)
- At least one seed in every “window” - guarantee
- Parameter w (window size) controls density

ch1 CACGACTCTGGTACCTAGACTCGATCGATCGTACTGT....
ch2 GCACACTGGTACCTAGACACGATCGCCCCGTTGTGTCA....
ch3 GCTGACTGTTACAACCTATACTCGATCGCCCCGTTGTCT....
...

?



AGACCCGAT...

Read (query)

Index

k-mer	Positions
CACG	[(ch1,0), (ch17,1202), ...]
AGAC	[(ch1,16), (ch2, 14), ...]
TATA	[(ch2, 25), (ch13, 205), ...]
GACT	...

...

1. Subsampling seeds

- Store only a subset of k-mers (aka *sketching* or *thinning*)
- At least one seed in every “window” - guarantee
- Parameter w (window size) controls density

ch1 CACGACTCTGGTACCTAGACTCGATCGATCGTACTGT....
ch2 GCACACTGGTACCTAGACACGATCGCCCCGTTGTGTCA....
ch3 GCTGACTGTTACAACCTATACTCGATCGCCCCGTTGTCT....
...

?
AGACCCGAT...
Read (query)

Index

k-mer	Positions
CACG	[(ch1,0), (ch17,1202), ...]
AGAC	[(ch1,16), (ch2, 14), ...]
TATA	[(ch2, 25), (ch13, 205), ...]
GACT	...
...	

Only minimizers in read are queried! → No hits!

1. Subsampling seeds

- Store only a subset of k-mers (aka *sketching* or *thinning*)
- At least one seed in every “window” - guarantee
- Parameter w (window size) controls density

ch1 CACGACTCTGGTACCTAGACTCGATCGATCGTACTGT....
ch2 GCACACTGGTACCTAGACACGATCGCCCCGTTGTGTCA....
ch3 GCTGACTGTTACAACCTATACTCGATCGCCCCGTTGTCT....
...

?
AGACCGAT...

Read (query)

Index

k-mer	Positions
CACG	: [(ch1,0), (ch17,1202), ...]
AGAC	: [(ch1,16), (ch2, 14), ...]
TATA	: [(ch2, 25), (ch13, 205), ...]
GACT	: ...
...	

Pick a smaller w to get more seeds

From now on: No more visually pleasant lexicographical order on minimizers!



In practice: A *hash function* is used to scramble ordering -> sample k-mers more uniformly across the set of k-mers

2. Chaining (of seeds)

- Find collinear “chains” of seeds
- Collinearity: same order in read and ref
- (Additional: Not too differing in distance)

2. Chaining (of seeds)

- Find collinear “chains” of seeds
- Collinearity: same order in read and ref
- (Additional: Not too differing in distance)


ch1 CACGACTCTGGTACCTAGACTCGATCGATCGTACTGT....
ch2 GCACACTGGTACCTAGACACGATCGCCCGTTGTGTCA....
ch3 GCTGACTGTTACAACATACTCGATCGCCCGTTGTCT....
...

AGACCCGATATCGTACTTGT

Read (query)


2. Chaining (of seeds)

- Find collinear “chains” of seeds
- Collinearity: same order in read and ref
- (Additional: Not too differing in distance)



2. Chaining (of seeds)

- Find collinear “chains” of seeds
- Collinearity: same order in read and ref
- (Additional: Not too differing in distance)




3. Extension

- Grow out base-level pairwise *extension alignment* from seeds
- Get an alignment score


3. Extension

- Grow out base-level pairwise *extension alignment* from seeds
- Get an alignment score



3. Extension

- Grow out base-level pairwise *extension alignment* from seeds
- Get an alignment score



3. Extension

- Grow out base-level pairwise *extension alignment* from seeds
- Get an alignment score

ch1 CACGACTCTGGTACCTAGACTCGATCGATCGTACT-GT...
 ||||| ||||| ||||| ||||| ||
Read (query) AGACCCGAT--ATCGTACTTTGT

Alignment score (AS): $Am - Bx - \sum_i O \cdot Eg_i$

A : Match score (2)

m : #matches

B : Mismatch penalty (4)

x : #mismatches

O : Gap open penalty (4)

g_i : length of gap i

E : Gap extension penalty (2)

3. Extension

- Grow out base-level pairwise *extension alignment* from seeds
- Get an alignment score

ch1 CACGACTCTGGTACCTAGACTCGATCGATCGTACT-GT...
 ||||| ||||| ||||| ||||| ||
Read (query) AGACCCGAT--ATCGTACTTGT

$$\text{Alignment score (AS): } Am - Bx - \sum_i O \cdot Eg_i$$

A : Match score (2)

m : #matches

B : Mismatch penalty (4)

x : #mismatches

O : Gap open penalty (4)


g_i : length of gap i

E : Gap extension penalty (2)

$$AS = 2 \cdot 18 - 4 \cdot 1 - (4 + 2 \cdot 2) - (4 + 2 \cdot 1) = 18$$

3. Extension

- Grow out base-level pairwise *extension alignment* from seeds
- Get an alignment score



3. Extension

- Grow out base-level pairwise *extension alignment* from seeds
- Get an alignment score

ch2 GCACACTGGTACCTAGACACGATCGCCCGTTGTGTCA....

Read (query) AGACCCGATATCGTACTTGT

3. Extension

- Grow out base-level pairwise *extension alignment* from seeds
- Get an alignment score

ch2 GCACACTGGTACCTAGACACGAT--CGCCCGTTGTGTCA....
 ||||| ||||| || | |||||
 AGACCCGATATCGTAC-TTGT

Alignment score (AS): $Am - Bx - \sum_i O \cdot Eg_i$

A: Match score (2)

m : #matches

B: Mismatch penalty (4)

x : #mismatches

O: Gap open penalty (4)

g_i : length of gap i

E: Gap extension penalty (2)

$AS = ?$

Long-read splice alignment

- Still seed-chain-extend, but:
 - Several collinear chains (think one per exon)
 - ‘Local’ extension alignments around exon sites
 - Splice site specific extension penalty (canonical site?)
- For small introns (same local chain): gap or intron penalty

Long-read splice alignment

- Still seed-chain-extend, but:
 - Several collinear chains (think one per exon)
 - ‘Local’ extension alignments around exon sites
 - Splice site specific extension penalty (canonical site?)
- For small introns (same local chain): gap or intron penalty

Reference

CACGACTCTGGTACCTAGACTCGATCGATCGTACTGTCGTATGCTA...

Read (query)

CGACTCTGGTACCGATCGATCTGTCGTATG

Long-read splice alignment


- Still seed-chain-extend, but:
 - Several collinear chains (think one per exon)
 - ‘Local’ extension alignments around exon sites
 - Splice site specific extension penalty (canonical site?)
- For small introns (same local chain): gap or intron penalty

Reference

CACGACTCTGGTACCTAGACTCGTTCCGATCGTACTGTCCGTATGCTA...

Read (query)

CGACTCTGGTACCGTTCCGATCTGTCCGTATG



Parameters (Minimap2)

Many parameters!

(and many 'hidden' parameters <https://lh3.github.io/minimap2/minimap2.html>)

- Seeding:
 - -k: size of seeds
 - -w: density of seeds
 - -f: fraction filtering repetitive of seeds
- Chaining:
 - -n Minimum seeds to include in a chain
 - -m minimum chaining score
- Alignment:
 - -a: extension alignment
 - Is the base level alignment off? Parameters -A -B -O -E
- For splice alignment:
 - The ensemble '-ax splice' flag typically take care of suitable parameter choices (canonical splice sites)
 - -u: how to find GT-AG
 - -G Maximum gap on the reference (long introns)
 - -C Cost for a non-canonical GT-AG splicing (Hidden)

Parameters (Minimap2)

Many parameters!

(and many 'hidden' parameters <https://lh3.github.io/minimap2/minimap2.html>)

Poor alignments?

- Error rate/sequence diversity
- Short exons
- Your reference characteristics (e.g., repetitiveness)
- Forgot 'cleaning' (trimming) reads before alignment
 - pychopper (ONT)
 - lima (PacBio)

Challenge: Get as many exons correctly aligned as possible with minimap2!

Prerequisites

- Install minimap2
- Install Python
- Download data:
 - `git clone https://github.com/ksahlin/misc.git`
 - (or download 'evaluate_instance.py' and the three files in dataset1 folder at <https://github.com/ksahlin/misc/tree/main/LongTREC/examples>)

Data

- 100kbp simulated 'genome' - nucleotides A,C,G,T *randomly simulated*
- Exons of size 1nt, 2nt, 3nt,... 99nt simulated.
- Challenge: Map an error-free transcript containing all exons

annotation.gtf								
genome	sim	gene	1	100000	.	+	.	gene_id "ENSG00000223972";
genome	sim	transcript	1	100000	.	+	.	gene_id "ENSG00000223972";
genome	sim	exon	1001	1001	.	+	.	gene_id "ENSG00000223972";
genome	sim	exon	2001	2002	.	+	.	gene_id "ENSG00000223972";
genome	sim	exon	3001	3003	.	+	.	gene_id "ENSG00000223972";
genome	sim	exon	4001	4004	.	+	.	gene_id "ENSG00000223972";
genome	sim	exon	5001	5005	.	+	.	gene_id "ENSG00000223972";
genome	sim	exon	6001	6006	.	+	.	gene_id "ENSG00000223972";
genome	sim	exon	7001	7007	.	+	.	gene_id "ENSG00000223972";
genome	sim	exon	8001	8008	.	+	.	gene_id "ENSG00000223972";
genome	sim	exon	9001	9009	.	+	.	gene_id "ENSG00000223972";
genome	sim	exon	10001	10010	.	+	.	gene_id "ENSG00000223972";
genome	sim	exon	11001	11011	.	+	.	gene_id "ENSG00000223972";

Challenge: Get as many exons correctly aligned as possible with minimap2!

SETUP

1. `git clone https://github.com/ksahlin/misc.git`
2. `cd misc/LongTREC/examples/`
3. `minimap2 --eqx -a [PARAMS] dataset1/genome.fa dataset1/query.fa > mm2.sam`
4. `python evaluate_instance.py --gtf dataset1/annotation.gtf --samfile mm2.sam`

Challenge: Get as many exons correctly aligned as possible with minimap2!

SETUP

1. `git clone https://github.com/ksahlin/misc.git`
2. `cd misc/LongTREC/examples/`
3. `minimap2 --eqx -a [PARAMS] dataset1/genome.fa dataset1/query.fa > mm2.sam`
4. `python evaluate_instance.py --gtf dataset1/annotation.gtf --samfile mm2.sam`

Try steps 3 and 4 on **dataset1** with various parameters

Challenge: Get as many exons correctly aligned as possible with minimap2!

SETUP

1. `git clone https://github.com/ksahlin/misc.git`
2. `cd misc/LongTREC/examples/`
3. `minimap2 --eqx -a [PARAMS] dataset1/genome.fa dataset1/query.fa > mm2.sam`
4. `python evaluate_instance.py --gtf dataset1/annotation.gtf --samfile mm2.sam`

Try steps 3 and 4 on **dataset1** with various parameters

OUTPUT

#Exact	#Approx	min-E	min-A
7	25	68	37

Challenge: Get as many exons correctly aligned as possible with minimap2!

SETUP

1. `git clone https://github.com/ksahlin/misc.git`
2. `cd misc/LongTREC/examples/`
3. `minimap2 --eqx -a [PARAMS] dataset1/genome.fa dataset1/query.fa > mm2.sam`
4. `python evaluate_instance.py --gtf dataset1/annotation.gtf --samfile mm2.sam`

Try steps 3 and 4 on **dataset1** with various parameters

OUTPUT

#Exact	#Approx	min-E	min-A
7	25	68	37

#Exact: Number of exact exon alignments

#Approx: Number of exact OR approximate exon alignments
(alnmt overlaps true exon)

min-E: Smallest exon with an exact exon alignment

min-A: Smallest exon with an exact OR approximate exon alignment nr_exact

Challenge: My attempts

Parameters

--eqx -a

--eqx -ax splice

#E

1

6

#A

1

70

E_min

99

43

A_min

99

18

Challenge: My attempts

Parameters

--eqx -a

--eqx -ax splice

--eqx -ax splice -k 10

--eqx -ax splice -k 10 -w 1

#E

#A

E_min

A_min

1

1

99

99

6

70

43

18

6

72

43

11

7

69

38

9

Challenge: My attempts

Parameters	#E	#A	E_min	A_min
--eqx -a	1	1	99	99
--eqx -ax splice	6	70	43	18
--eqx -ax splice -k 10	6	72	43	11
--eqx -ax splice -k 10 -w 1	7	69	38	9
--eqx -ax splice -u n	41	79	18	18
--eqx -ax splice -u n -k 10	42	81	16	11
--eqx -ax splice -u n -k 10 -w 1	40	77	32	9

Challenge: My attempts

Parameters	#E	#A	E_min	A_min
--eqx -a	1	1	99	99
--eqx -ax splice	6	70	43	18
--eqx -ax splice -k 10	6	72	43	11
--eqx -ax splice -k 10 -w 1	7	69	38	9
--eqx -ax splice -u n	41	79	18	18
--eqx -ax splice -u n -k 10	42	81	16	11
--eqx -ax splice -u n -k 10 -w 1	40	77	32	9
--eqx -ax splice -u n -k 10 -B 10 -O 4,12	53	89	12	11
--eqx -ax splice -u n -k 10 -w 1 -B 10 -O 4,12	53	91	12	9

Challenge: My attempts

Parameters	#E	#A	E_min	A_min
--eqx -a	1	1	99	99
--eqx -ax splice	6	70	43	18
--eqx -ax splice -k 10	6	72	43	11
--eqx -ax splice -k 10 -w 1	7	69	38	9
--eqx -ax splice -u n	41	79	18	18
--eqx -ax splice -u n -k 10	42	81	16	11
--eqx -ax splice -u n -k 10 -w 1	40	77	32	9
--eqx -ax splice -u n -k 10 -B 10 -O 4,12	53	89	12	11
--eqx -ax splice -u n -k 10 -w 1 -B 10 -O 4,12	53	91	12	9
--eqx -ax splice -k 7 -w 1 -u n -B 10 -O 4,12	24	42	12	11

Challenge: My attempts

Parameters	#E	#A	E_min	A_min
--eqx -a	1	1	99	99
--eqx -ax splice	6	70	43	18
--eqx -ax splice -k 10	6	72	43	11
--eqx -ax splice -k 10 -w 1	7	69	38	9
--eqx -ax splice -u n	41	79	18	18
--eqx -ax splice -u n -k 10	42	81	16	11
--eqx -ax splice -u n -k 10 -w 1	40	77	32	9
--eqx -ax splice -u n -k 10 -B 10 -O 4,12	53	89	12	11
--eqx -ax splice -u n -k 10 -w 1 -B 10 -O 4,12	53	91	12	9
--eqx -ax splice -k 7 -w 1 -u n -B 10 -O 4,12	24	42	12	11

- Difficult to predict final outcome due to all heuristics and parameters (-w 1)
- Useful to know the basics: k-mer size, fitting to canonical splice sites.
- Extension alignment parameters (A, B, O, E) remains guesswork to this day.
 - But: lower O, E penalties prefers to open gaps - helpful to splice alignments

Three additional datasets

- Dataset2: 7% errors
- Dataset3: No errors, only GT-AG
- Dataset4: 7% errors, only GT-AG

Three additional datasets

- **Dataset2: 7% errors**
- Dataset3: No errors, only GT-AG
- Dataset4: 7% errors, only GT-AG

Parameters	#E	#A	E_min	A_min
--eqx -a	1	1	93	93
--eqx -ax splice	5	57	61	22
--eqx -ax splice -k 10	5	64	61	17
--eqx -ax splice -k 10 -w 1	5	66	61	11
--eqx -ax splice -u n	29	73	26	22
--eqx -ax splice -u n -k 10	30	73	25	17
--eqx -ax splice -u n -k 10 -w 1	29	76	26	11
--eqx -ax splice -u n -k 10 -B 10 -O 4,12	33	83	17	17
--eqx -ax splice -u n -k 10 -w 1 -B 10 -O 4,12	34	89	16	11
--eqx -ax splice -k 7 -w 1 -u n -B 10 -O 4,12	12	32	16	15

- With errors/mutations - k and w is more influential
- Less exact matches - but almost the same amount of exons with approximate mappings

Three additional datasets

- Dataset2: 7% errors
- **Dataset3: No errors, only GT-AG**
- Dataset4: 7% errors, only GT-AG

Parameters	#E	#A	E_min	A_min
--eqx -a	0	1	-	98
--eqx -ax splice	76	81	21	16
--eqx -ax splice -k 10	76	82	21	12
--eqx -ax splice -k 10 -w 1	73	78	24	9
--eqx -ax splice -u n	38	80	25	16
--eqx -ax splice -u n -k 10	38	81	25	12
--eqx -ax splice -u n -k 10 -w 1	38	78	25	9
--eqx -ax splice -u n -k 10 -B 10 -O 4,12	45	88	12	12
--eqx -ax splice -u n -k 10 -w 1 -B 10 -O 4,12	48	91	9	9
--eqx -ax splice -k 7 -w 1 -u n -B 10 -O 4,12	18	41	10	10

- Minimap2 is optimised for canonical splice sites!
- However, tuned alignment parameters still finds the most exon sites (91)

Three additional datasets

- Dataset2: 7% errors
- Dataset3: No errors, only GT-AG
- **Dataset4: 7% errors, only GT-AG**

Parameters	#E	#A	E_min	A_min
--eqx -a	0	1	-	99
--eqx -ax splice	63	71	32	14
--eqx -ax splice -k 10	61	72	35	13
--eqx -ax splice -k 10 -w 1	63	76	32	10
--eqx -ax splice -u n	23	72	32	14
--eqx -ax splice -u n -k 10	23	76	32	13
--eqx -ax splice -u n -k 10 -w 1	23	75	32	10
--eqx -ax splice -u n -k 10 -B 10 -O 4,12	33	84	16	14
--eqx -ax splice -u n -k 10 -w 1 -B 10 -O 4,12	34	88	12	10
--eqx -ax splice -k 7 -w 1 -u n -B 10 -O 4,12	9	29	12	10

- With errors/mutations - k and w is more influential
- Minimap2 is optimised for canonical splice sites!
- However, tuned alignment parameters still finds the most exon sites (88)

Take-home of the analysis

minimap2

- ✓ Minimap2 ‘just works’ (installation and running easy compared to competition)
- ✓ There is likely a parameter combination that is suitable for your needs
- The problem is how to find it
- Minimap2 works better with canonical splice sites

General splice alignment


- Don’t blindly trust the input alignments for your isoform detection
- Know your dataset and genome
- If you are interested in smaller exons - there are better tools (deSALT, uLTRA)

Alternative aligners: uLTRA

- + uLTRA runs minimap2 under-the-hood to find alignments in unannotated regions
- + It picks the best alignment (alignment score) between uLTRA and minimap2

Alternative aligners: uLTRA

Preparation



Alternative aligners: uLTRA


Preparation

Exon
annotations

Parts

Flanks

Segments



Alternative aligners: uLTRA

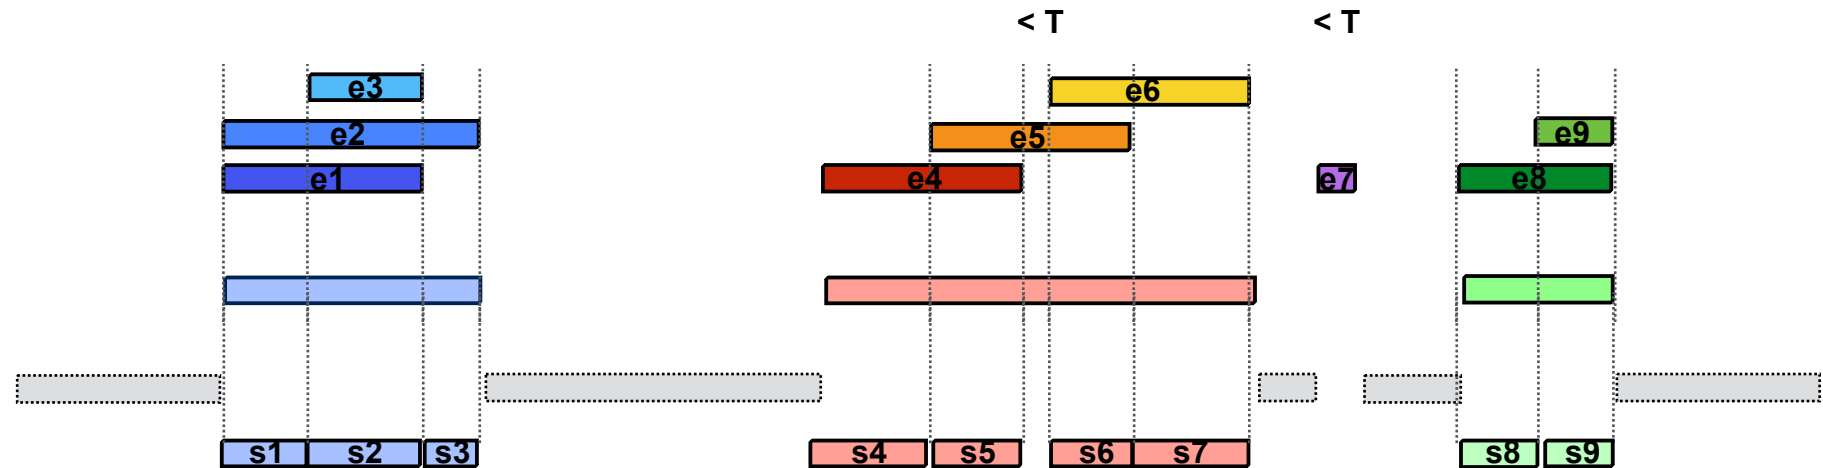
Preparation

Exon
annotations

Parts

Flanks

Segments



Alignment

MEM chaining
to parts and flanks

Read 1



Read 2



Alternative aligners: uLTRA


Preparation

Exon
annotations

Parts

Flanks

Segments



Alignment

Read 1 MEMs



MEM chaining
to parts and flanks

Read 1



Read 2



Alternative aligners: uLTRA


Preparation

Exon
annotations

Parts

Flanks

Segments




Alignment

MEM chaining
to parts and flanks

Read 1 MEMs

Read 1

Read 2



Alternative aligners: uLTRA


Preparation

Exon
annotations

Parts

Flanks

Segments




Alignment

MEM chaining
to parts and flanks

Read 1 MEMs

Read 1



Read 2



Alternative aligners: uLTRA


Preparation

Exon
annotations

Parts

Flanks


Segments



Alignment

Read 1 MEMs

Read 2 MEMs



MEM chaining
to parts and flanks

Read 1



Read 2



Alternative aligners: uLTRA

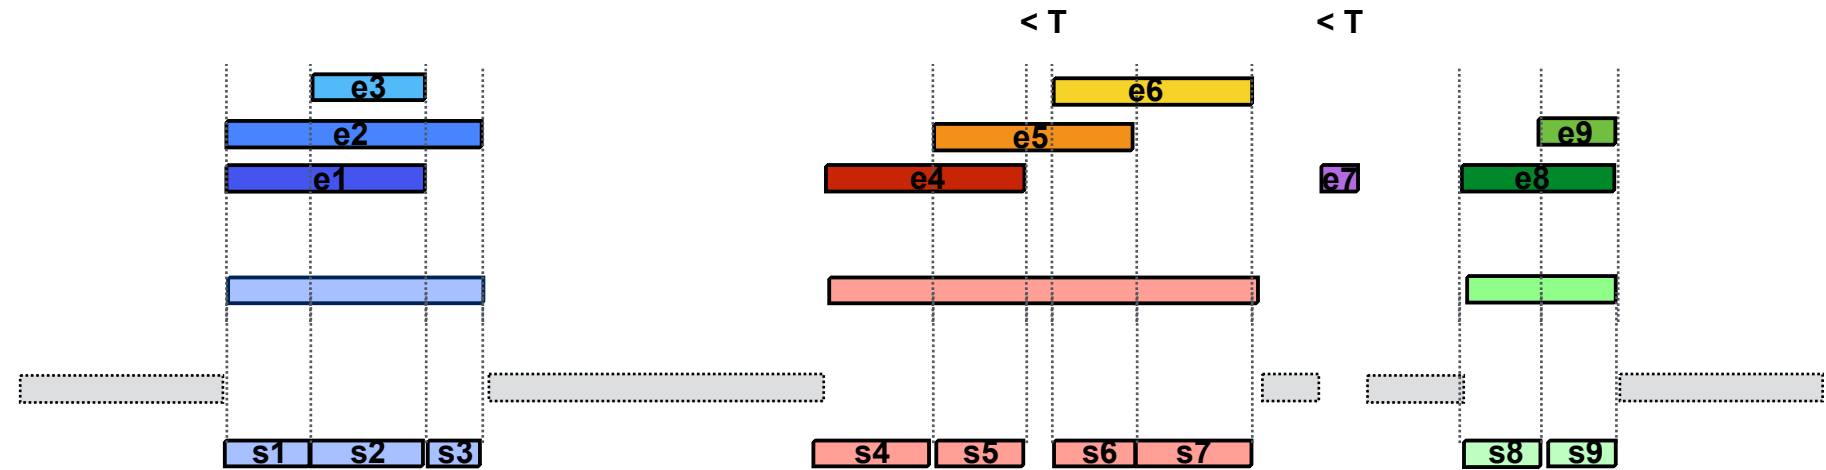
Preparation

Exon annotations

Parts

Flanks

Segments



Alignment

Read 1 MEMs

Read 2 MEMs



MEM chaining to parts and flanks

Read 1





Read 2



Segment chaining

* Small exons and exons containing small segments are included in MAM chaining

e5 e7



Alternative aligners: uLTRA


Preparation

Exon annotations

Parts

Flanks

Segments



Alignment


Read 1 MEMs

Read 2 MEMs



MEM chaining to parts and flanks

Read 1





Read 2




Segment chaining

* Small exons and exons containing small segments are included in MAM chaining


e5 e7



Read 1



Read 2



Reference

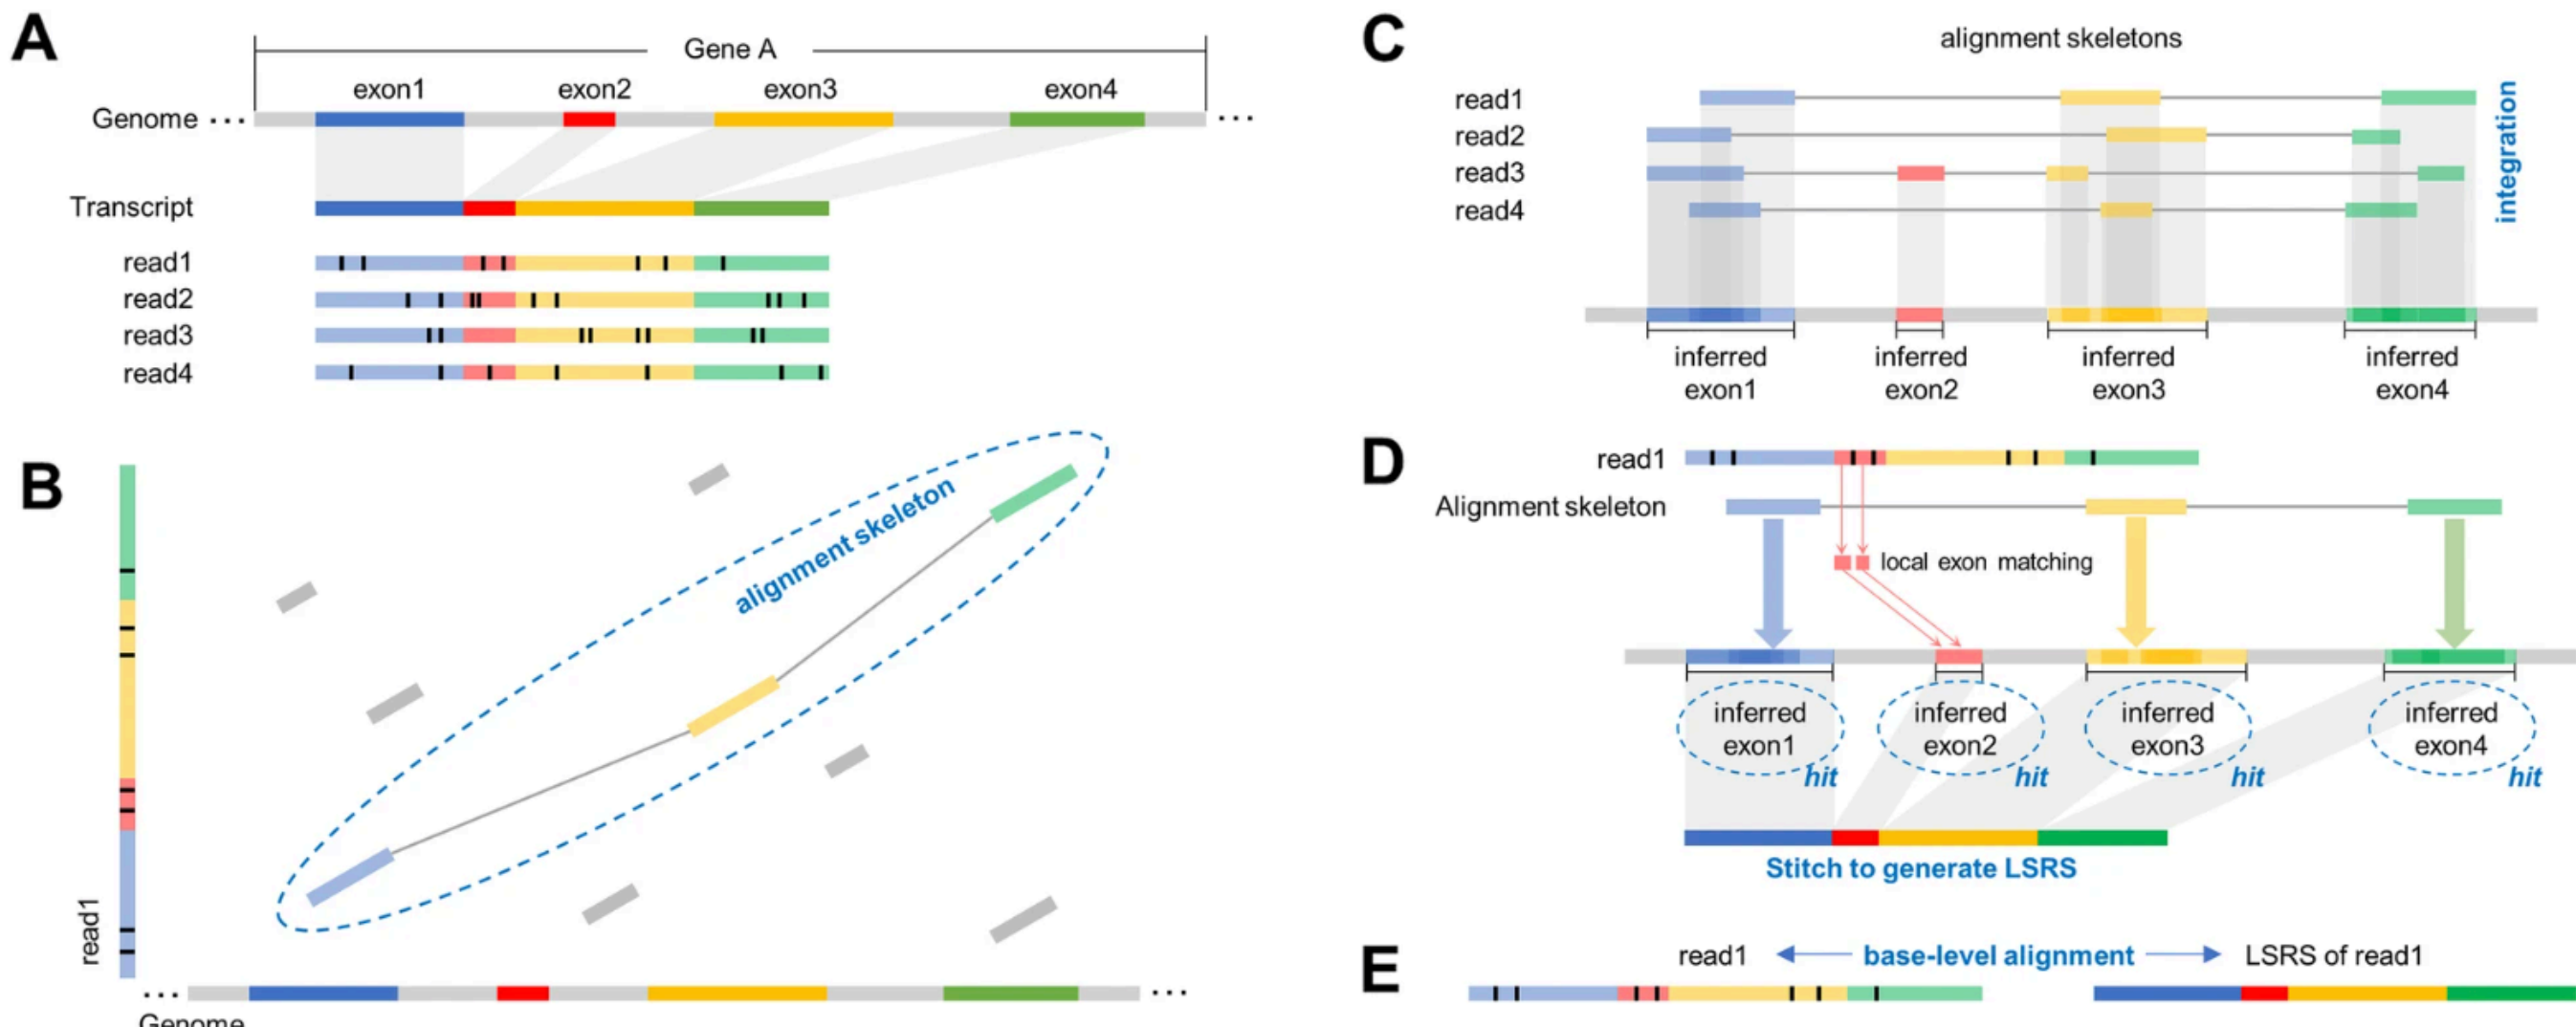
uLTRA v0.0.4.1 results

	Parameters	#E	#A	E_min	A_min
Dataset 1	default	94	95	5	5
Dataset 2	default	94	95	5	5
Dataset 4	default	94	95	5	5
Dataset 3	default	93	94	5	5

Take-home uLTRA

- ✓ Uses annotation to refine alignments
- ✓ Robust against different datasets (error rate, canonical or not..)
- ✓ Good for finding small exons (that are part of the annotation)
- ✓ Has minimap2's accuracy in unannotated regions
- Slower than minimap2 (at least 2 times slower)
- Disk space: prints a lot of intermediate/temporary files

Alternative aligners: deSALT





- Could not install on my MacBooks (Intel and M1 chip)
- *Prediction based on previous analysis*: places somewhere between minimap2 and uLTRA in accuracy on these datasets

Accuracy binned by exon size (uLTRA, deSALT, minimap2)

Data from uLTRA paper

No errors (mapping transcripts to genome)

Transcript annotations with 7-8% errors



Interpreting the output (SAM)

The alignment of reads:

```
ref      AGCATGTTAGATAA**GATAGCTGTGCTAGTAGGCAGTCAGCGCCAT
+r002     aaaAGATAA*GGATA
```


Interpreting the output (SAM)

The alignment of reads:

```
ref      AGCATGTTAGATAA**GATAGCTGTGCTAGTAGGCAGTCAGCGCCAT
+r002     aaaAGATAA*GGATA
```

The resulting SAM file:

```
@HD VN:1.6 SO:coordinate
```

```
@SQ SN:ref LN:45
```

```
r002    0 ref    9 30 3S6M1P1I4M *    0    0 AAAAGATAAGGATA    *
```

Interpreting the output (SAM)

The alignment of reads:

```
ref      AGCATGTTAGATAA**GATAGCTGTGCTAGTAGGCAGTCAGCGCCAT
+r002    aaaAGATAA*GGATA
```

The resulting SAM file:

```
@HD VN:1.6 SO:coordinate
@SQ SN:ref LN:45
r002    0 ref 9 30 3S6M1P1I4M * 0 0 AAAAGATAAGGATA *
```

A 'flag' Ref start MAPQ CIGAR



Interpreting the output (SAM)

The alignment of reads:

```
ref      AGCATGTTAGATAA**GATAGCTGTGCTAGTAGGCAGTCAGCGCCAT
+r002     aaaAGATAA*GGATA
+r003     gcctaAGCTAA
```

The resulting SAM file:

```
@HD VN:1.6 SO:coordinate
@SQ SN:ref LN:45
r002    0 ref 9 30 3S6M1P1I4M * 0 0 AAAAGATAAGGATA *
r003    0 ref 9 30 5S6M * 0 0 GCCTAAGCTAA * SA:Z:ref,29,-,6H5M,17,0;
```

Diagram illustrating the mapping of SAM file fields to their biological meaning:

- A 'flag' (points to the 0 in the first read's flag field)
- Ref start (points to the 9 in the first read's reference start field)
- MAPQ (points to the 0 in the first read's mapping quality field)
- CIGAR (points to the 3S6M1P1I4M in the first read's CIGAR field)

Interpreting the output (SAM)

The alignment of reads:

```

ref      AGCATGTTAGATAA**GATAGCTGTGCTAGTAGGCAGTCAGCGCCAT
+r002      aaaAGATAA*GGATA
+r003      gcctaAGCTAA
+r004      ATAGCT.....TCAGC
  
```

The resulting SAM file:

```

@HD VN:1.6 SO:coordinate
@SQ SN:ref LN:45
r002    0 ref  9 30 3S6M1P1I4M * 0 0 AAAAGATAAGGATA *
r003    0 ref  9 30 5S6M      * 0 0 GCCTAAGCTAA   * SA:Z:ref,29,-,6H5M,17,0;
r004    0 ref 16 30 6M14N5M   * 0 0 ATAGCTTCAGC   *
  
```

Annotations for the SAM file:

- A 'flag'**: Points to the first field (0) of the alignment line.
- Ref start**: Points to the reference start coordinate (9).
- MAPQ**: Points to the mapping quality score (0).
- CIGAR**: Points to the CIGAR string (3S6M1P1I4M).

Interpreting the output (SAM)

The alignment of reads:

```

ref      AGCATGTTAGATAA**GATAGCTGTGCTAGTAGGCAGTCAGCGCCAT
+r002      aaaAGATAA*GGATA
+r003      gcctaAGCTAA
+r004      ATAGCT.....TCAGC
-r003      ttagctTAGGC
  
```

The resulting SAM file:

```

@HD VN:1.6 SO:coordinate
@SQ SN:ref LN:45
r002    0 ref  9 30 3S6M1P1I4M * 0 0 AAAAGATAAGGATA *
r003    0 ref  9 30 5S6M      * 0 0 GCCTAAGCTAA   * SA:Z:ref,29,-,6H5M,17,0;
r004    0 ref 16 30 6M14N5M   * 0 0 ATAGCTTCAGC   *
r003 2064 ref 29 17 6H5M      * 0 0 TAGGC        * SA:Z:ref.9+.5S6M.30.1:
  
```

Annotations for the SAM file:

- A 'flag'**: Points to the value 2064 in the SAM record for r003.
- Ref start**: Points to the reference start coordinate 9 in the SAM record for r002.
- MAPQ**: Points to the mapping quality value 0 in the SAM record for r002.
- CIGAR**: Points to the CIGAR string 3S6M1P1I4M in the SAM record for r002.

Interpreting the output (SAM)

The alignment of reads:

```

ref      AGCATGTTAGATAA**GATAGCTGTGCTAGTAGGCAGTCAGCGCCAT
+r002      aaaAGATAA*GGATA
+r003      gcctaAGCTAA
+r004      ATAGCT.....TCAGC
-r003      ttagctTAGGC

```

The resulting SAM file:

```

@HD VN:1.6 SO:coordinate
@SQ SN:ref LN:45
r002    0 ref  9 30 3S6M1P1I4M * 0 0 AAAAGATAAGGATA *
r003    0 ref  9 30 5S6M      * 0 0 GCCTAAGCTAA   * SA:Z:ref,29,-,6H5M,17,0;
r004    0 ref 16 30 6M14N5M  * 0 0 ATAGCTTCAGC   *
r003 2064 ref 29 17 6H5M    * 0 0 TAGGC        * SA:Z:ref.9.+.5S6M.30.1:

```

Annotations with arrows:

- A 'flag' (points to 2064)
- Ref start (points to 29)
- MAPQ (points to 0)
- CIGAR (points to 6H5M)

Useful links:

- SAM format: <https://samtools.github.io/hts-specs/SAMv1.pdf>
- Explain SAM flags: <https://broadinstitute.github.io/picard/explain-flags.html>

Trouble shooting and visualisation

- **Samtools:**
 - quick sanity check statistics (% aligned reads etc)
- **BLAT** - server version (<https://genome.ucsc.edu/cgi-bin/hgBlat>)
 - align a single or handful of reads with the absolute best accuracy
 - Extremely slow
 - Not the latest references
 - good for checking selected reads
- **Seaview** (<https://doua.prabi.fr/software/seaview>)
 - Align reads against the selves - transcript redundancy redundancy checking etc
- **IGV** (<https://igv.org/doc/desktop/>):
 - Stacked read view - *can sometimes be* good for gene/transcript level visualisation (SNPs etc)

References

- **Seed-chain-extend:** Sahlin, K., Baudeau, T., Cazaux, B. *et al.* A survey of mapping algorithms in the long-reads era. *Genome Biol* **24**, 133 (2023). <https://doi.org/10.1186/s13059-023-02972-3>
- **uLTRA:** Kristoffer Sahlin, Veli Mäkinen, Accurate spliced alignment of long RNA sequencing reads, *Bioinformatics*, Volume 37, Issue 24, December 2021, Pages 4643–4651, <https://doi.org/10.1093/bioinformatics/btab540>
- **Minimap2:** Li H. (2018) Minimap2: pairwise alignment for nucleotide sequences. *Bioinformatics*, 34, 3094–3100.
- **deSALT:** Liu, B., Liu, Y., Li, J. et al. deSALT: fast and accurate long transcriptomic read alignment with de Bruijn graph-based index. *Genome Biol* 20, 274 (2019) doi:10.1186/s13059-019-1895-9
- **SAM format:** Heng Li, Bob Handsaker, Alec Wysoker, Tim Fennell, Jue Ruan, Nils Homer, Gabor Marth, Goncalo Abecasis, Richard Durbin, 1000 Genome Project Data Processing Subgroup, The Sequence Alignment/Map format and SAMtools, *Bioinformatics*, Volume 25, Issue 16, August 2009, Pages 2078–2079, <https://doi.org/10.1093/bioinformatics/btp352>
- **Samtools:** Heng Li, Bob Handsaker, Alec Wysoker, Tim Fennell, Jue Ruan, Nils Homer, Gabor Marth, Goncalo Abecasis, Richard Durbin, 1000 Genome Project Data Processing Subgroup, The Sequence Alignment/Map format and SAMtools, *Bioinformatics*, Volume 25, Issue 16, August 2009, Pages 2078–2079, <https://doi.org/10.1093/bioinformatics/btp352>
- **BLAT:** <https://genome.ucsc.edu/cgi-bin/hgBlat>